



Scala the Cloud Native Way: Lessons Learned from Two Years of Linkerd in Production

Dennis Adjei-Baah

A little about me



Software Engineer @ **Buoyant**

Linkerd contributor

Twitter: @dadjuib

Github: @dadjiebaah



Linkerd

- Service mesh
 - HTTP, HTTP/2, gRPC request routing
 - reliable, secure, visible
- Built on Finagle/Netty
- Runs as an app with your microservices.



Why you would use it?

- Service discovery integration
- Load balancing
 - Request level load balancing
- Retries
- Advanced routing
 - Canary deployments
 - Blue/Green deployments
- TLS

Where it all started: 2015



Introducing linkerd

linkerd is a dynamic linker for distributed applications (aka "microservices"). In the same way that ``ld(1)`` binds software components (libraries), linkerd binds services by mediating inter-service communication (RPC).

linkerd builds upon finagle & netty--Twitter's JVM networking stack--and it exposes many of the advanced operational features developed by Twitter, Soundcloud, and other large internet applications.



24+ months in production

2k+ Slack channel members

7,000+ GitHub stars

20m+ DockerHub pulls

80+ contributors

400b+ production requests/mo




credit karma



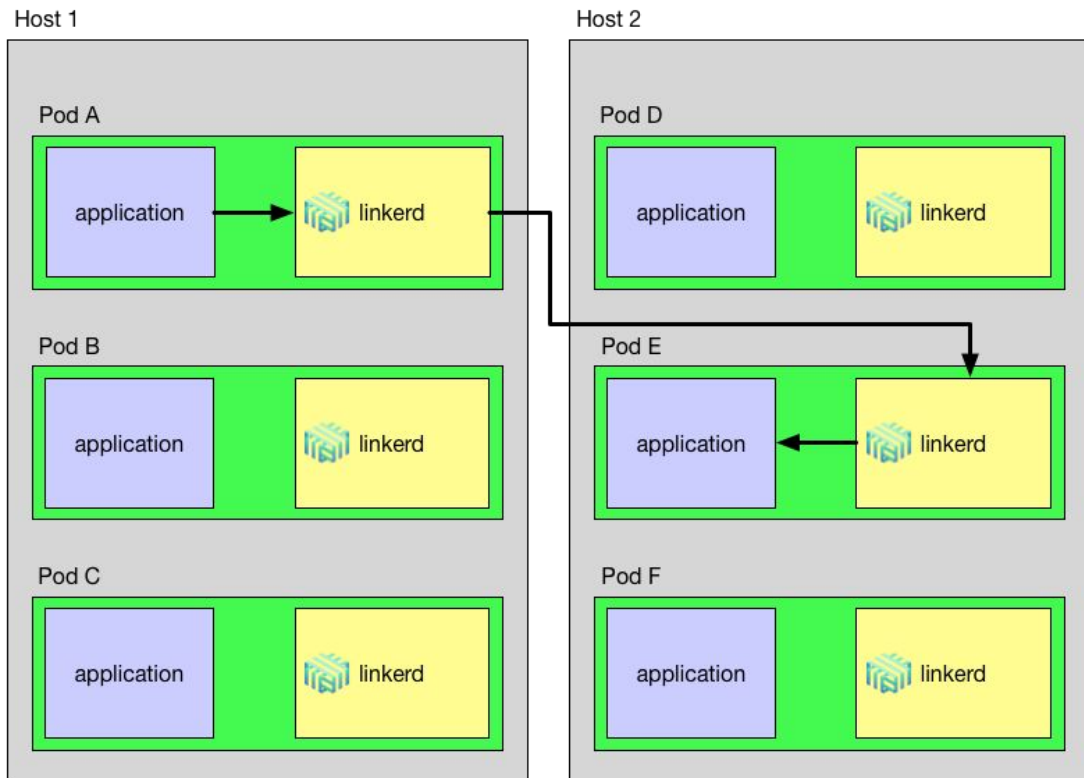
Challenges over those three years

1. Memory usage
2. Memory usage
3. Memory usage

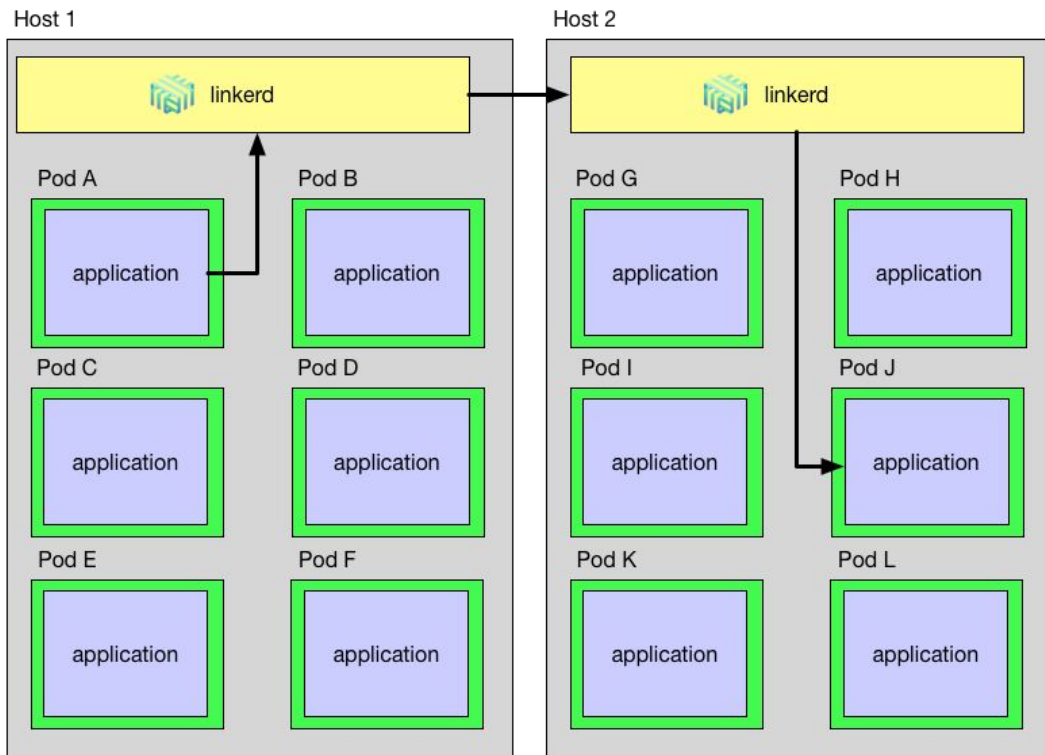
The world is going cloud native




Linkerd runs as a distributed proxy



Alternative deployment: Per Host



Linkerd in production today



Optimizing Linkerd in production

Ways to reduce JVM memory footprint

- JVM tuning
- Use OpenJ9
- Use GraalVM
- Don't use the JVM

Tuning JVM, Finagle, Netty

- JVM flags
 - Concurrent mark sweep collector, tiered compilation many more
- Finagle
 - Number of threads used to handle concurrency
 - Amount of memory used per thread to tame memory footprint
- Netty
 - Has similar settings help control memory used for


Finding the right tuning parameters for all environments is hard!

OpenJ9

- Built by IBM Java development team
- Tested for decades
- Designed for the cloud




OpenJ9 benchmarks




Linkerd + OpenJ9

Before



After




Other OpenJ9 optimizations

- **-Xshareclasses**
 - Store class data to disk to increase subsequent start up time
- **-XX:+IdleTuningGcOnIdle**
 - Trigger GC when VM perceives application as idle
- **-Xtune:virtualized**
 - Signals the VM that it running in a constrained cloud environment which makes it use threads more efficiently
- **-Xgcpolicy**
 - Configure GC behavior that optimizes garbage collection for a specific type of application.

Learn more [here](#)

Linkerd + OpenJ9 tuned



GraalVM

- One VM for multiple programming languages (R, Python, Java and more)
- Ahead of time compilation
- Build native images from Scala



Work ongoing at Oracle

- Finagle and Netty on GraalVM
 - With some code modifications
- Making progress on getting Linkerd to work.
- Linkerd's plugin system

```
private[linkerd] lazy val LoadedInitializers = Initializers(  
  LoadService[ProtocolInitializer],  
  LoadService[NamerInitializer],  
  LoadService[InterpreterInitializer] :+ DefaultInterpreterInitializer,  
  LoadService[TransformerInitializer],  
  LoadService[IdentifierInitializer],  
  LoadService[ResponseClassifierInitializer],  
  LoadService[TelemeterInitializer],  
  LoadService[AnnouncerInitializer],  
  LoadService[FailureAccrualInitializer],  
  LoadService[RequestAuthorizerInitializer],  
  LoadService[TracePropagatorInitializer]  
)
```

Contribute to Linkerd+GraalVM working group

- Working group mailing list
 - <https://lists.cncf.io/g/cncf-linkerd-graal-wg>
- Linkerd GraalVM slack channel
 - <https://linkerd.slack.com/messages/CB3S9LMV5>

Don't use the JVM

- Only if you have really strict memory and throughput goals!
- JVM is more than capable in supporting high throughput
- Linkerd must be transparent
 - Require low system requests
 - Low latency added to proxied request



Current state of Linkerd

- Officially supported: OpenJDK with tuning
- Experimental: OpenJ9
- Future: GraalVM
- Minimalist rewrite: Linkerd 2.x
 - Off the JVM
 - Very few features compared to Linkerd 1.x (For now)

Recap

- Cloud native poses some challenges when working with the JVM
- JVM has a rich ecosystem e.g. we can build on top of Finagle and Netty
- Moving off the JVM is expensive
- End of the day all implementation details
- What problems we solve with Linkerd

Stay connected

Upcoming events:

- December 10: [Linkerd 101](#) hands on class at KubeCon + CNCFCon Seattle
- December 11-13: [lots of Linkerd talks](#) at KubeCon!

Community:

- Join the [linkerd-users mailing list](#)
- Join the [Linkerd Slack](#) (2300+ users)
- **Twitter:** @dadjeb
- **Github:** @dajeibaah (Slides will be available there)